

Applied Algorithm Design: Exam

Prof. Pietro Michiardi

Rules and suggestions

- The idea is to complete first the questions and then the exercises.
- Having **the good** answer for all questions will guarantee to pass the exam with 10/20.
- Exercises are difficult if you are not familiar with the course contents. Suggestion: pick one exercise that you think is more feasible for you and focus on that one.

Questions

1. Given a digraph $G = (V, E)$ and two nodes s and t , find the minimum number of edges whose removal disconnects t from s .
 - (a) Given this problem, which algorithm can be used to find the solution?
 - (b) Which is the running time of this algorithm?
 - (c) Why is this question is important in real systems?
2. Given an undirected graph $G = (V, E)$, name a simple algorithm to check for bipartiteness.
3. Draw a simple (e.g. 5 nodes) graph, and illustrate (in a sequence of graphs) at least two iterations of the PageRank algorithm: in particular, each node in the drawings should indicate what is the current PageRank value.
4. In your home works, you had to implement the Fair Sojourn Protocol (FSP).
 - What is the motivation for FSP? Why FIFO or processor sharing are not enough?
 - What is the fundamental assumption of FSP, and by the way for any size-based scheduling protocol?

- In your opinion, what is the impact of wrong estimates on the job sizes?

5. Why is it important for the Gale-Shapley algorithm to output the same result no matter what the proposal order is? Be brief.

Exercise 1

You work for a security company that needs to obtain licenses for n different pieces of cryptographic software. Due to regulations, you can only obtain those licenses at the rate of at most one per month.

Each license is currently selling for a price of \$100. However, they are all becoming more expensive according to exponential growth curves: in particular, the cost of license j increases by a factor of $r_j > 1$ each month, where r_j is a given parameter. This means that if license j is purchased t months from now, it will cost $100 \cdot r_j^t$. We will assume that all the price growth rates are distinct; that is, $r_i \neq r_j$ for licenses $i \neq j$ (even though they start at the same price \$100).

The question is: Given that the company can only buy at most one license a month, in which order should it buy the licenses so that the total amount of money it spends is as small as possible?

Give an algorithm that takes the n rates of price growth r_1, r_2, \dots, r_n and computes an order in which to buy the licenses so that the total amount of money spent is minimized. **The running time of your algorithm should be polynomial in n .**

[Hint] Here we look for a greedy algorithm. Now, since you have to prove optimality, it is helpful to recall that when a greedy algorithm works optimally by putting a set of things in an optimal order, it is often effective to try proving correctness using an **exchange argument**.]

Exercise 2

You're doing some stress-testing on various models of glass jars to determine the height from which they can be dropped and still not break. The setup for this experiment, on a particular type of jar, is as follows. You have a ladder with n rungs, and you want to find the highest rung from which you can drop a copy of the jar and not have it break. We call this the *highest safe rung*.

It might be natural to try binary search: drop a jar from the middle rung, see if it breaks, and then recursively try from rung $n/4$ or $3n/4$ depending on the outcome. But this has the drawback that you could break a lot of jars in finding the answer.

If your primary goal were to conserve jars, on the other hand, you could try the following strategy. Start by dropping a jar from the first rung, then the second rung, and so forth, climbing one higher each time until the jar breaks. In this way, you only

need a single jar – at the moment it breaks, you have the correct answer – but you may have to drop it n times (rather than $\log n$ as in the binary search solution).

So here is the trade-off: it seems you can perform fewer drops if you're willing to break more jars. To understand better how this tradeoff works at a quantitative level, let's consider how to run this experiment given a fixed "budget" of $k \geq 1$ jars. In other words, you have to determine the correct answer – the highest safe rung – and can use at most k jars in doing so.

- (a) Suppose you are given a budget of $k = 2$ jars. Describe a strategy for finding the highest safe rung that requires you to drop a jar at most $f(n)$ times, for some function $f(n)$ that grows slower than linearly.

[Hint] It should be the case that $\lim_{n \rightarrow \infty} f(n)/n = 0$.

- (b) Now suppose you have a budget of $k > 2$ jars, for some given k . Describe a strategy for finding the highest safe rung using at most k jars. If $f_k(n)$ denotes the number of times you need to drop a jar according to your strategy, then the functions f_1, f_2, f_3, \dots should have the property that each grows asymptotically slower than the previous one.

[Hint] $\lim_{n \rightarrow \infty} f_k(n)/f_{k-1}(n) = 0$ for each k .