

# Applied Algorithm Design

## Lecture 7

Pietro Michiardi

Eurecom

# Randomized Algorithms

- Randomization and probabilistic analysis are themes that cut across many areas of CS, including algorithm design
- There are two main ways of looking at randomized algorithms:
  - 1 **Average analysis:** in this case we consider traditional algorithms confronted with random input. Here we study the behavior of an algorithm on an “average” input
  - 2 **Randomized algorithms:** this is an approach where randomization is purely internal to the algorithm, which takes random decisions along its execution

## Why should we look at algorithms that make random decisions?

- By allowing randomization, we make the underlying model more powerful
- Efficient **deterministic algorithms** that always yield the correct answer are a special case of efficient randomized algorithms, that only need to yield the correct answer **with high probability** [a.k.a. Monte Carlo algorithms]
- Deterministic algorithms are also special cases of randomized algorithms that are always correct, and run efficiently **in expectation** [a.k.a. Las Vegas algorithms]

- We will look at randomized algorithms for a number of problems where there exist comparably efficient deterministic algorithms
  - ▶ A randomized algorithm may be conceptually much simpler
  - ▶ A randomized algorithm may function while maintaining very little internal state or memory of the past
- For **distributed systems**, in which many loosely interacting processes operate, randomized algorithms can reduce the amount of explicit communication or synchronization, and can be seen as a tool for **symmetry-breaking** among processes

Do you need to be an expert in probability theory to design and analyze randomized algorithms?

- Of course, knowledge only helps, but very little tools from probability theory are enough to analyze a wide range of algorithms
- In this lecture, we will revise some important concepts that are used all along the analysis of randomized algorithms, such as Union Bounds, and Chernoff Bounds

# Applications of Randomized Algorithms

# Content Resolution

- We begin with a first application of randomized algorithms: content resolution in a distributed system
- This is a typical (simple) example of the general style of analysis we will use for the analysis of randomized algorithms
- We will work on notions related to **events**, their **independence** and use a simple **Union Bound**



## The problem

- We have  $n$  processes  $P_1, P_2, \dots, P_n$ , each competing for access to a **single shared** database
- Time is slotted into discrete **rounds**
- The database can be accessed by at most one process in a single round
  - ▶ If two processes attempt a simultaneous access, they will both be “locked out”
  - ▶ If no process access the database in a round, then the round will be “lost”

# Content Resolution

## Observations

- Each process wants to access the database as often as possible
  - It is pointless for all processes to try to concurrently access the database at all rounds, for they would be all “locked out”
- ⇒ We need a way to divide up the rounds among the processes in an **equitable** way

## Communication

- If processes can communicate with one another, then it is possible to find many ways of addressing this problem
- We will instead assume that **no communication** is allowed between processes to coordinate

# Content Resolution

## A randomized algorithm

- Randomization provides a natural protocol for this problem
- For some number  $p > 0$  that we'll determine shortly, each process will attempt to access that database in each round with probability  $p$ , **independently** of the decisions of the other processes
  - ⇒ If exactly one process decides to access in a given round, it will succeed
  - ⇒ If two or more try, then they will collide
  - ⇒ If none try, then the round is wasted

This is a symmetry-breaking paradigm, used to “smooth out” contention

## Analyzing the algorithm

- We start by defining some basic events and think about their probabilities
- For a given process  $P_i$  and a given round  $t$ , let  $\mathcal{A}[i, t]$  denote the event that  $P_i$  attempts to access the database at time  $t$
- We know that each process attempts with a probability  $p$  in every round, so  $\Pr[\mathcal{A}[i, t]] = p, \forall t > 0$
- For every event, there is a **complementary event**,  $\overline{\mathcal{A}[i, t]}$ , indicating that  $P_i$  does not attempt to access the database in round  $t$ :

$$\overline{\mathcal{A}[i, t]} = 1 - \Pr[\mathcal{A}[i, t]] = 1 - p$$

Our real concern is whether a process **succeeds** in accessing the DB in a given round

- Denote the success event by  $\mathcal{S}[i, t]$
- Clearly,  $P_i$  must attempt an access in round  $t$  in order to succeed  
 $\Rightarrow P_i$  access the DB at time  $t$  and each other process **does not** attempt an access to the DB at time  $t$
- Thus,  $\mathcal{S}[i, t]$  is equal to the intersection of the event  $\mathcal{A}[i, t]$  with all the complementary events  $\overline{\mathcal{A}[j, t]}$ ,  $\forall j \neq i$ :

$$\mathcal{S}[i, t] = \mathcal{A}[i, t] \cap \left( \bigcap_{j \neq i} \overline{\mathcal{A}[j, t]} \right)$$

By assumption, access attempts are **independent**

- By the independency of the events, the intersection corresponds to the product of events probabilities:

$$\Pr [\mathcal{S}[i, t]] = \Pr [\mathcal{A}[i, t]] \cdot \prod_{j \neq i} \Pr [\overline{\mathcal{A}[j, t]}] = p(1 - p)^{n-1}$$

- This is a **closed-form expression** for the probability that  $P_i$  succeeds in accessing the DB in round  $t$

### How can we maximize the success probability?

- We will play on selecting a good value for  $p$ 
  - ▶ If  $p = 0$  or  $p = 1$ , then  $\Pr[\mathcal{S}[i, t]] = 0$
- Let's examine the function  $f(p) = p(1 - p)^{n-1}$ 
  - ▶  $f(p)$  is positive for  $p \in (0, 1)$
  - ▶

$$f'(p) = (1 - p)^{n-1} - (n - 1)p(1 - p)^{n-2}$$

which has a single zero at the value  $p = 1/n$ , where the maximum is achieved

Intuitively, this corresponds to having a single process to make an attempt in each distinct round, which also guarantees some sort of “fairness”

## Observations

- When we set  $p = 1/n$  we have:

$$\Pr[S[i, t]] = \frac{1}{n} \left(1 - \frac{1}{n}\right)^{n-1}$$

## Reminder from basic calculus:

- The function  $\left(1 - \frac{1}{n}\right)^n$  converges monotonically from  $\frac{1}{4}$  up to  $\frac{1}{e}$  as  $n$  increases from 2
- The function  $\left(1 - \frac{1}{n}\right)^{n-1}$  converges monotonically from  $\frac{1}{2}$  up to  $\frac{1}{e}$  as  $n$  increases from 2



## Asymptotic analysis

- It is useful getting a sense of the asymptotic value of  $\Pr[\mathcal{S}[i, t]]$
- From basic calculus we have:

$$\frac{1}{en} \leq \Pr[\mathcal{S}[i, t]] \leq \frac{1}{2n}$$

$$\Rightarrow \Pr[\mathcal{S}[i, t]] = \Theta(1/n)$$

### How long until a process succeeds?

- Assume the algorithm we designed, using the optimal probability  $p = 1/n$
- We can see from the previous inequality that the probability of a process  $P_i$  to succeed in any one round  $t$  is not very high, especially if  $n$  is reasonably large
- How about considering **multiple rounds**?

### How long until a process succeeds?

- Let  $\mathcal{F}[i, t]$  denote the “failure event” that process  $P_i$  does not succeed in **any** of the rounds 1 through  $t$
- This is clearly the intersection of the complementary events  $\overline{\Pr[S[i, r]]}$  for  $r = 1, 2, \dots, t$
- Since each of these events is **independent**, we have:

$$\Pr[\mathcal{F}[i, t]] = \Pr\left[\bigcap_{r=1}^t \overline{\Pr[S[i, r]]}\right] = \prod_{r=1}^t \overline{\Pr[S[i, r]]} =$$
$$\left[1 - \frac{1}{n} \left(1 - \frac{1}{n}\right)^{n-1}\right]^t$$

Let's think asymptotically again, to avoid complicated expressions

- Recall that the probability of success was  $\Theta(1/n)$  after one round
- Specifically, it was bounded between  $1/(en)$  and  $1/(2n)$
- Using the expression above, we have:

$$\Pr[\mathcal{F}[i, t]] = \prod_{r=1}^t \overline{\Pr[\mathcal{S}[i, r]]} \leq \left(1 - \frac{1}{en}\right)^t$$

We can still simplify by substitution

- Let  $t = \lceil en \rceil$ , then we have:

$$\Pr \left[ \mathcal{F}[i, t] \right] \leq \left( 1 - \frac{1}{en} \right)^{\lceil en \rceil} \leq \left( 1 - \frac{1}{en} \right)^{en} \leq \frac{1}{e}$$

This is a very compact and useful asymptotic statement: the probability that process  $P_i$  does not succeed in any of rounds 1 to  $\lceil en \rceil$  is upper-bounded by the constant  $e^{-1}$ , independently of  $n$

## Content Resolution

How can we manipulate time such as the failure probability is very small?

- Let's set  $t = \lceil en \rceil (c \cdot \ln n)$ , then we have:

$$\Pr [\mathcal{F}[i, t]] \leq \left(1 - \frac{1}{en}\right)^t \leq \left(\left(1 - \frac{1}{en}\right)^{\lceil en \rceil}\right)^{c \ln n} \leq e^{-c \ln n} = n^{-c}$$

Asymptotically, we can say that:

- After  $\Theta(n)$  rounds, the probability that  $P_i$  has not yet succeeded is bounded by a constant
- Between then and  $\Theta(n \ln n)$  rounds, this probability drops a lot, bounded by an inverse polynomial in  $n$

Ok, so we know everything about one process, but the question is: how much time to wait before **all** processes get through?

- We say that the protocol **fails** after  $t$  rounds if some process has not yet succeeded in accessing the DB
  - Let  $\mathcal{F}_t$  denote the event that the protocol fails after  $t$  rounds
- ⇒ Our goal is to find a reasonably small  $t$  such as  $\mathcal{F}_t$  is small

- The event  $\mathcal{F}_t$  occurs if and only if one of the events  $\mathcal{F}[i, t]$  occurs, which writes as:

$$\mathcal{F}_t = \bigcup_{i=1}^n \mathcal{F}[i, t]$$

### ATTENTION:

We have a union of events that are **not independent**

- We will use the **Union Bound**, which says that the probability of a union of events is upper-bounded by the sum of their individual probabilities



## The Union Bound

Given events  $\mathcal{E}_1, \mathcal{E}_2, \dots, \mathcal{E}_n$  we have:

$$\Pr \left[ \bigcup_{i=1}^n \mathcal{E}_i \right] \leq \sum_{i=1}^n \Pr [\mathcal{E}_i]$$

- Note that this is not an equality
- The bound is good enough when the union represents a “bad event” that we are trying to avoid, and we want a bound on its probability

- We have that:

$$\mathcal{F}_t = \bigcup_{i=1}^n \mathcal{F}[i, t]$$

$$\Pr[\mathcal{F}_t] \leq \sum_{i=1}^n \Pr[\mathcal{F}[i, t]]$$

- The expression on the right hand-side is a sum of  $n$  terms with the same value
- To make the probability of  $\mathcal{F}_t$  small, we need to make each of the terms on the right hand-side significantly smaller than  $1/n$

- We know that choosing  $t = \Theta(n)$  will not be good enough
- If we choose  $t = \lceil en \rceil \cdot c \ln n$  then we have that  $\Pr[\mathcal{F}[i, t]] \leq n^{-c}, \forall i$
- Precisely, we can set  $t = 2\lceil en \rceil \ln n$ , which gives us:

$$\Pr[\mathcal{F}_t] \leq \sum_{i=1}^n \Pr[\mathcal{F}[i, t]] \leq n \cdot n^{-2} = n^{-1}$$

## Theorem:

With probability at least  $1 - n^{-1}$ , all processes succeed in accessing the DB at least once within  $t = 2 \lceil en \rceil \ln n$  rounds

# Finding the Global Minimum Cut

## The problem

- Given an undirected graph  $G = (V, E)$ , we define a **cut** of  $G$  to be a partition of  $V$  into two non-empty sets  $A$  and  $B$ 
  - ▶ Note the difference with an  $s - t$  cut: here we don't have a source nor a sink for flow, and actually we didn't even need to define a flow
- For a cut  $(A, B)$  in an undirected graph  $G$ , the **size** of  $(A, B)$  is the number of edges with one end in  $A$  and the other in  $B$
- A **global minimum cut** is a cut of minimum size

## Why is it useful to study cuts?

The global min-cut is a **robustness** parameter: it is the smallest number of edges whose deletion disconnects the graph

# Finding the Global Minimum Cut

## Proposition

There is a polynomial-time algorithm to find a global min-cut in an undirected graph  $G$

### Proof.

- We use the similarity with  $s - t$  cuts in directed graphs
- We transform a given graph  $G$  by replacing every **undirected** edge with two oppositely oriented directed edges, and call it  $G'$
- We pick two arbitrary nodes  $s, t \in V$  and find the  $s - t$  min-cut in  $G'$
- If  $(A, B)$  is a min-cut in  $G'$  it is also so in  $G$
- We repeat the procedure for every other node as a sink, setting  $t \in V \setminus \{s\}$
- Hence, we need to compute  $n - 1$   $s - t$  min-cuts, and the best among them will be the **global min-cut**



# Finding the Global Minimum Cut


## Observation

- It looks like the global min-cut is harder than what we have seen before
- In reality, it turns out that we don't really need to compute all the  $n - 1$  min-cuts
- This can be shown with some difficult tricks, and we'll not do it here
- Also the following randomized algorithm, that is called the **Contraction Algorithm**, can benefit from some tricks and run definitively faster than what we show here

## Finding the Global Minimum Cut

Designing the algorithm: Contraction algorithm [Karger 1995]

- Pick an edge  $e = (u, v)$  uniformly at random
- **Contract** edge  $e$ 
  - ▶ replace  $u$  and  $v$  by single new super-node  $w$
  - ▶ preserve edges, updating endpoints of  $u$  and  $v$  to  $w$
  - ▶ keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes  $v_1$  and  $v_2$
- Return the cut (all nodes that were contracted to form  $v_1$ )





# Finding the Global Minimum Cut

## Analyzing the algorithm

- The algorithm is making random choices
- ⇒ there is some probability that it will succeed in finding a global min-cut, and some probability that it won't
- One might think that the success probability is exponentially small: there are exponentially many possible cuts of  $G$
- In reality, the success probability is polynomially small, which means that we can run the algorithm many times (polynomial number of times) and return the best cut found so far
- ⇒ With high probability, we will find a global min-cut

## Finding the Global Minimum Cut

### Theorem:

The Contraction Algorithm returns a global min-cut of  $G$  with probability at least  $\frac{1}{\binom{n}{2}}$

- We focus on a global min-cut  $(A, B)$  of  $G$  and suppose it has size  $k$   
 $\Rightarrow$  There is a set  $F$  of  $k$  edges with one end in  $A$  and the other in  $B$

We want to give a lower bound on the probability that the Contraction Algorithm returns the cut  $(A, B)$

## Finding the Global Minimum Cut

- Consider what could go wrong with the algorithm
  - ▶ What if an edge in  $F$  were contracted?
  - ▶ A node of  $A$  and a node of  $B$  would get thrown together in a super-node, and  $(A, B)$  could not be returned by the algorithm
  - ▶ Conversely, if an edge not in  $F$  is contracted, then there is still a chance that  $(A, B)$  could be returned

We want an upper bound on the probability that an edge in  $F$  is contracted

## Finding the Global Minimum Cut

- We need first a lower bound on the cardinality of  $E$
  - Note that if any node  $v$  had degree less than  $k$ , then the cut  $(\{v\}, V \setminus \{v\})$  would have size less than  $k$
- ⇒ this would contradict our assumption that  $(A, B)$  is a global min-cut
- Hence, every node in  $G$  has degree at least  $k$ , and so:

$$|E| \geq \frac{1}{2}kn$$

- ⇒ The probability that an edge in  $F$  is contracted is at most:

$$\frac{k}{\frac{1}{2}kn} = \frac{2}{n}$$

## Finding the Global Minimum Cut

- Now, consider the situation after  $j$  iterations, when there are  $n - j$  super-nodes in the current graph  $G'$
  - Assume that no edge in  $F$  has been contracted so far
  - Every cut in  $G'$  is a cut of  $G$  and so there are at least  $k$  edges incident to every super-node of  $G'$
- ⇒  $G'$  has at least  $\frac{1}{2}k(n - j)$  edges, and the probability that an edge of  $F$  is contracted in the next iteration  $j + 1$  is at most:

$$\frac{k}{\frac{1}{2}k(n - j)} = \frac{2}{n - j}$$

## Finding the Global Minimum Cut

The cut  $(A, B)$  will be returned by the algorithm if no edge of  $F$  is contracted in any of iterations  $1, 2, \dots, n-2$

- Let  $\mathcal{E}_j$  be the event that an edge of  $F$  is not contracted in iteration  $j$   
 $\Rightarrow$  We have that

$$\Pr[\mathcal{E}_1] \geq 1 - \frac{2}{n}$$

$$\Pr[\mathcal{E}_{j+1} | \mathcal{E}_1 \cap \mathcal{E}_2 \dots \cap \mathcal{E}_j] \geq 1 - \frac{2}{n-j}$$

We are interested in lower bounding the quantity:

$$\Pr[\mathcal{E}_1 \cap \mathcal{E}_2 \dots \cap \mathcal{E}_{n-2}]$$

## Finding the Global Minimum Cut

- If we unwind the formula for conditional probability we have:

$$\begin{aligned}\Pr[\mathcal{E}_1 \cap \mathcal{E}_2 \dots \cap \mathcal{E}_{n-2}] &= \\&= \Pr[\mathcal{E}_1] \cdot \Pr[\mathcal{E}_2|\mathcal{E}_1] \cdots \Pr[\mathcal{E}_{j+1}|\mathcal{E}_1 \cap \mathcal{E}_2 \cdots \cap \mathcal{E}_j] \\&\quad \cdots \Pr[\mathcal{E}_{n-2}|\mathcal{E}_1 \cap \mathcal{E}_2 \cdots \cap \mathcal{E}_{n-3}] \\&\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{n-j}\right) \cdots \left(1 - \frac{2}{3}\right) \\&= \left(\frac{n-2}{2}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \cdots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\&= \frac{2}{n(n-1)} = \binom{n}{2}^{-1}\end{aligned}$$

## Finding the Global Minimum Cut

So what if we want to make this probability very small?

Repeat the algorithm many times! But how many times is enough?

- A single run of the Contraction Algorithm fails to find a global min-cut with probability at most  $1 - 1/\binom{n}{2}$ 
  - ▶ This number is very close to 1
  - ▶ Need to **amplify** our probability of success
- If we run the algorithm  $\binom{n}{2}$  times, then the probability to fail is at most:

$$\left(1 - \frac{1}{\binom{n}{2}}\right)^{\binom{n}{2}} \leq \frac{1}{e}$$

- If we run the algorithm  $\binom{n}{2} \ln n$  times, then we have that the probability of failure is at most:

$$e^{-\ln n} = \frac{1}{n}$$



# Finding the Global Minimum Cut

## Remark

- Overall running time is slow since we perform  $\Theta(n^2 \ln n)$  iterations and each takes  $\Omega(m)$  time
- Improvement: [Karger-Stein 1996]  $O(n^2 \log^3 n)$ 
  - ▶ Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when  $n/\sqrt{2}$  nodes remain
  - ▶ Run contraction algorithm until  $n/\sqrt{2}$  nodes remain
  - ▶ Run contraction algorithm **twice** on resulting graph, and return best of two cuts
- Extensions: Naturally generalizes to handle positive weights

## Best known Contraction Algorithm [Karger 2000]

Runs in  $O(m \ln^3 n)$ , which is faster than best known max flow algorithm or deterministic global min cut algorithm

# Background

# Random Variables and Their Expectations

- Thus far our analysis of randomized algorithms and processes has been based on identifying certain “bad events”, and bounding their probabilities
  - Here we want to look at a quantitative style of analysis by considering certain parameters related to a randomized algorithms, such as its running time or the quality of the produced solution
- ⇒ We seek to determine the **expected** size of these parameters over the random choices made by the algorithm

# Random Variables and Their Expectations

- **Expectation.** Given a discrete random variable  $X$ , its expectation  $E[X]$  is defined as:

$$E[X] = \sum_{j=0}^{\infty} j \Pr[X = j]$$

## Random Variables and Their Expectations

- **Waiting for a first success.** Coin is heads with probability  $p$  and tails with probability  $1 - p$ . How many independent flips  $X$  until first heads?

$$\begin{aligned} E[X] &= \sum_{j=0}^{\infty} j \Pr[X = j] = \sum_{j=0}^{\infty} j(1-p)^{j-1}p = \\ &= \frac{p}{1-p} \sum_{j=0}^{\infty} j(1-p)^j = \frac{p}{1-p} \cdot \frac{1-p}{p^2} = \frac{1}{p} \end{aligned}$$

This was a more useful example, in which we see how an appropriate random variable lets us talk about something like the “running time” of a simple random process

## Random Variables and Their Expectations

- **Useful property.** If  $X$  is a boolean random variable, then:

$$E[X] = \Pr[X = 1]$$

$$E[X] = \sum_{j=0}^{\infty} j \cdot \Pr[X = j] = \sum_{j=0}^1 j \cdot \Pr[X = j] = \Pr[X = 1]$$

- **Linearity of expectation.** Given two random variables  $X$  and  $Y$ , not necessarily independent, defined over the same probability space, we have that:

$$E[X + Y] = E[X] + E[Y]$$

This last fact **decouples** a complex calculation into simpler pieces

## Random Variables and Their Expectations

- **Game.** Shuffle a deck of  $n$  cards; turn them over one at a time; try to guess each card
- **Memoryless guessing.** No psychic abilities; can't even remember what's been turned over already. Guess a card from full deck uniformly at random
- **Claim.** The expected number of correct guesses is 1

### Using linearity of expectation

- Let  $X_i = 1$  if the  $i^{th}$  prediction is correct and 0 otherwise
- Let  $X$  be the number of correct guesses  $X = X_1 + X_2 + \dots + X_n$

$$E[X_i] = \Pr[X_i = 1] = 1/n$$

$$E[X] = E[X_1 + \dots + X_n] = E[X_1] + \dots + E[X_n] = \sum_{i=1}^n \frac{1}{n} = 1$$

## Random Variables and Their Expectations

- **Game.** Shuffle a deck of  $n$  cards; turn them over one at a time; try to guess each card
- **Guessing with memory.** Guess a card uniformly at random from cards not yet seen
- **Claim.** The expected number of correct guesses is  $\Theta(\ln n)$

### Proof: (Using linearity of expectation)

- Let  $X_i = 1$  if  $i$ -th prediction is correct, and 0 otherwise
- Let  $X = X_1 + \dots + X_n$  be the number of correct guesses
- $E[X_i] = \Pr[X_i = 1] = 1/(n - i + 1)$
- $E[X] = E[X_1] + \dots + E[X_n] = 1/n + \dots + 1/2 + 1/1 = \sum_{i=1}^n \frac{1}{i} = H(n)$
- Where  $\ln(n+1) \leq H(n) \leq 1 + \ln n$  is the **Harmonic number**



## Random Variables and Their Expectations

- **Coupon collector.** Each box of cereal contains a coupon. There are  $n$  different types of coupons. Assuming all boxes are equally likely to contain each coupon, how many boxes before you have  $\geq 1$  coupon of each type?
- **Claim.** The expected number of steps is  $\Theta(n \ln n)$

### Proof:

- Phase  $j$ : time between  $j$  and  $j + 1$  distinct coupons
- Let  $X_j$  be the number of steps you spend in phase  $j$
- Let  $X = X_0 + \dots + X_{n-1}$  be the total number of steps

$$E[X] = \sum_{j=0}^{n-1} E[X_j] = \sum_{j=0}^{n-1} \frac{n}{n-j} = n \sum_{i=1}^n \frac{1}{i} = nH(n)$$

- Probability of success is  $\frac{n-j}{n} \Rightarrow$  expected waiting time  $\frac{n}{n-j}$

# Universal Hashing

- Randomization has also proved to be a powerful technique in the design of **data structures**
- Here we discuss a technique called **hashing**, which can be used to maintain a dynamically changing set of elements

## Applications:

- File systems
- DB
- P2P networks
- Web Caching

## The problem:

- Given a universe  $U$  of possible elements, maintain a subset  $S \subseteq U$  so that inserting, deleting, and searching in  $S$  is efficient
  - Note that  $S$  is generally a tiny fraction of  $U$
  - We need to create a **Dictionary**
- 
- Dictionary interface:
    - ▶ `Create()` : Initialize a dictionary with  $S = \phi$
    - ▶ `Insert( $u$ )` : Add element  $u \in U$  to  $S$
    - ▶ `Delete( $u$ )` : Delete  $u$  from  $S$ , if  $u$  is currently in  $S$
    - ▶ `Lookup( $u$ )` : Determine whether  $u$  is in  $S$

## Challenge:

Universe  $U$  can be extremely large so defining an array of size  $|U|$  is infeasible

## Note:

- We encountered already problems in which we were asked to maintain a list of dynamic elements: e.g. for BFS and DFS algorithms
- There the size of the set was known as an input to the algorithms, and it was feasible to maintain it with a traditional data structure
- Here we consider a case in which  $|U|$  is huge, hence it cannot fit in memory

# Universal Hashing

- **Hash function:**  $h : U \rightarrow 0, 1, \dots, n - 1$
- **Hashing.** Create an array  $H$  of size  $n$ . When processing element  $u$ , access array element  $H[h(u)]$

## Collision:

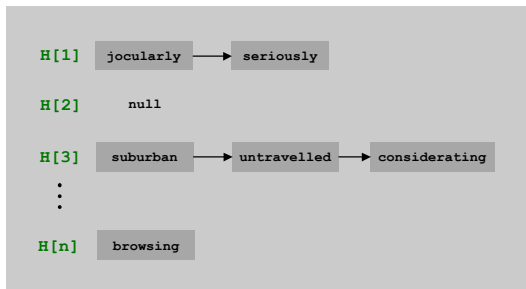
When  $h(u) = h(v)$  but  $u \neq v$

- A collision is expected after  $\Theta(\sqrt{n})$  random insertions
- This phenomenon is known as the **birthday paradox**

# Universal Hashing

## Separate chaining:

$H[i]$  stores linked list of elements  $u$  with  $h(u) = i$



# Universal Hashing

```
int h(String s, int n) {  
    int hash = 0;  
    for (int i = 0; i < s.length(); i++)  
        hash = (31 * hash) + s[i];  
    return hash % n;  
}
```

hash function ala Java string library

- **Deterministic hashing.** If  $|U| \geq n^2$ , then for any fixed hash function  $h$ , there is a subset  $S \subseteq U$  of  $n$  elements that all hash to same slot. Thus,  $\Theta(n)$  time per search in worst-case
- **Question.** Aren't ad hoc hash functions good enough in practice?

When can't we live with ad hoc hash function?

- **Denial-of-service attacks:** malicious adversary learns your ad hoc hash function (e.g., by reading Java API) and causes a big pile-up in a single slot that grinds performance to a halt
- Real world exploits. [Crosby-Wallach 2003]
  - ▶ Bro server: send carefully chosen packets to DOS the server, using less bandwidth than a dial-up modem
  - ▶ Perl 5.8.0: insert carefully chosen strings into associative array
  - ▶ Linux 2.4.20 kernel: save files with carefully chosen names



## Idealistic hash function

Maps  $m$  elements uniformly at random to  $n$  hash slots

- Running time depends on length of chains
- Average length of chain =  $\alpha = m/n$
- Choose  $n \sim m \Rightarrow$  on average  $O(1)$  per insert, lookup, or delete

- **Challenge.** Achieve idealized randomized guarantees, but with a hash function where you can easily find items where you put them
- **Approach.** Use randomization in the choice of  $h$

$\Rightarrow$  An adversary knows the randomized algorithm you're using, but doesn't know random choices that the algorithm makes

## Universal class of hash functions. [Carter-Wegman 1980s]

- For any pair of elements  $u, v \in U$ , we have  $\Pr_{h \in H} [h(u) = h(v)] \leq 1/n$ , where  $h$  is chosen uniformly at random
- Requirement 1: you must be able to select random  $h$  efficiently
- Requirement 2: you must be able to compute  $h(u)$  efficiently

# Universal Hashing

## Example:

$U = a, b, c, d, e, f, n = 2$

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1

$H = \{h_1, h_2\}$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1$$

$$\Pr_{h \in H} [h(a) = h(d)] = 0$$

...

not universal

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1
$h_3(x)$	0	0	1	0	1	1
$h_4(x)$	1	0	0	1	1	0

$H = \{h_1, h_2, h_3, h_4\}$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(d)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(e)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(f)] = 0$$

...

universal

# Universal Hashing

## Universal hashing property

- Let  $H$  be a universal class of hash functions
  - Let  $h \in H$  be chosen uniformly at random from  $H$
  - Let  $u \in U$
- ⇒ For any subset  $S \subseteq U$  of size at most  $n$ , the expected number of items in  $S$  that collide with  $u$  is at most 1

### Proof.

- For any element  $s \in S$ , define indicator random variable  $X_s = 1$  if  $h(s) = h(u)$  and 0 otherwise
- Let  $X$  be a random variable counting the total number of collisions with  $u$

$$E_{h \in H}[X] = E\left[\sum_{s \in S} X_s\right] = \sum_{s \in S} E[X_s] = \sum_{s \in S} \Pr[X_s = 1] \leq \sum_{s \in S} \frac{1}{n} = |S| \frac{1}{n} \leq 1$$

□

## Randomized On-line Caching

## Introduction

- We now discuss the use of randomization for the caching problem
- We begin by developing a class of algorithms: the **Marking Algorithms**
- We derive general performance guarantees on all marking algorithms and then focus on a randomized version

## The problem: **Cache Maintenance**

- We consider a processor whose full memory has  $n$  addresses
- It is equipped with a **cache** with  $k$  slots that can be accessed very quickly
- We can keep copies of  $k$  items from the full memory in cache slots, and when a memory location is accessed, the processor will first check the cache to see if it can be quickly retrieved

# Randomized Caching

## Useful definitions

- **Cache hit:** the cache contains the requested item
- **Cache miss:** the cache does not contain the item, and the processor needs to seek it from main memory
- **Cache eviction:** the action of eliminating one item from the cache

## Assumption:

We assume that the cache is kept full at all times



## Objective

The goal of a cache maintenance algorithm is to **minimize** the number of cache misses

- The sequence of memory references is not under the control of the algorithm
- ⇒ Which item currently in the cache should be evicted on each cache miss?

- **Off-line version:** always evict the item that will be needed **farthest in the future**. This is the optimal solution to the problem, which constitutes an **absolute benchmark**.
- It requires full knowledge of future items that will be needed

## Randomized Caching

- **On-line version:** in this case we cannot assume we have full knowledge of future requests
  - We can only base our eviction decisions on an history of past requests
- 
- In practice, the most commonly used eviction policy is the **LRU policy**
  - Least-recently-used: the intuition is that algorithms tend to have a certain **locality** in accessing data generally using the same set of data frequently and for a while. If a data item has not been accessed for a long time, this is a sign that it may not be accessed again for a long time

## A relation to approximation algorithms

- Here we will evaluate the performance of different eviction policy without making any assumptions (*i.e.*, locality) on the sequence of requests
- We will compare the number of misses made by an eviction policy on a sequence  $\sigma$  with the **minimum** number of misses it is possible to do on  $\sigma$
- We use  $f(\sigma)$  to denote the minimum number of misses, which is achieved by the farthest-in-future policy
- Comparing eviction policy to the optimum resembles to what we did with approximation algorithms
  - ▶ Then, we had a NP-hard problem and analyzed the gap obtained by polynomial time approximation algorithms
  - ▶ Here, we have an optimal algorithm! However, optimality requires full knowledge, but in practice we work on-line

## Designing a Class of Marking Algorithms

- The bounds on LRU and its randomized version will follow from a general template for designing **online** eviction policies
- To do well against the benchmark of  $f(\sigma)$ , we need an eviction policy that is sensitive to the difference between:
  - 1 In the recent past, the request sequence has contained more than  $k$  distinct items
  - 2 In the recent past, the requested sequence has come exclusively from a set of at most  $k$  items

# Randomized Caching

- In case 1 we know that  $f(\sigma)$  must be increasing, since no algorithm can handle more than  $k$  distinct items without incurring a cache miss
- In case 2 it is possible that  $\sigma$  is passing through a long stretch in which an optimal algorithm need not incur any misses at all

## Outline of the Marking Algorithm

- The algorithm prefers evicting items that don't seem to have been used in a long time
- The algorithm operates in **phases**
- In the following slide, we describe one phase of a Marking Algorithm

# Randomized Caching

---

## Algorithm 1: A general Marking Algorithm

---

Each memory item can be either **marked** or **unmarked**

At the beginning of the phase, all items are unmarked

On a request to item  $s$ :

Mark  $s$

**if**  $s$  *is in the cache* **then**

    | Evict nothing

**else**

    | **if** *All items currently in the cache are marked* **then**

        | Declare the phase over

        | Processing of  $s$  is deferred to start of next phase

    | **else**

        | **Evict** an **unmarked item** from the cache

    | **end**

**end**

---

# Randomized Caching

## Observations

- We have seen a class of algorithms, rather than a single specific algorithm
- The ambiguity comes from the **evict** and **unmarked item** emphasized before
  - ▶ How to evict and which unmarked item to select?
- Since a phase starts with all items unmarked, and items become marked only when accessed, the unmarked items have all been accessed less recently than the marked ones
- At any point in a phase, if there are any unmarked items in the cache, then the least recently used item must be unmarked

## Fact:

The LRU policy is a marking algorithm



## Analyzing Marking Algorithms

- Consider an arbitrary marking algorithm operating on a request sequence  $\sigma$
- We picture an optimal caching algorithm on  $\sigma$  alongside this marking algorithm, incurring an overall cost  $f(\sigma)$
- Assume there are  $r$  phases in this sequence  $\sigma$

## Padding

- We are going to “pad” the sequence  $\sigma$  both at the beginning and the end with some extra requests
- These will not add any extra misses to the optimal algorithm  $\Rightarrow$  the optimality bound on the padded sequence applies also to  $\sigma$

# Randomized Caching

## Phase 0

- Phase 0 takes place before the first phase  $\rightarrow$  all the items initially in the cache are requested once
- This does not affect the cost of either the marking algorithm or the optimal algorithm

## Final Phase

- Phase  $r$  ends with an epilogue in which every item currently in the cache of the optimal algorithm is requested twice in round-robin fashion
- This does not increase  $f(\sigma)$
- By the end of the second pass, the marking algorithm will contain each of the items in the cache, and each will be marked

## Performance bound

- We need two things:
  - 1 An upper bound on the number of misses incurred by the marking algorithm
  - 2 A lower bound saying that the optimum must incur at least a certain number of misses

## The history of a phase

- At the beginning of a phase, all items are unmarked
- Any item that is accessed during the phase is marked, and remains in the cache for the remainder of the phase
- Over the course of the phase, the number of marked items goes from 0 to  $k$ , and the next phase begins with a request to a  $(k + 1)$  item, different from all of these marked items

## Proposition:

- In each phase,  $\sigma$  contains accesses to exactly  $k$  distinct items
- The subsequent phase begins with an access to a different  $(k + 1)$  item

# Randomized Caching

## Proposition:

- The marking algorithm incurs at most  $k$  misses per phase, for a total of at most  $kr$  misses over all  $r$  phases

## Proposition:

- The optimum incurs at least  $r - 1$  misses

$$\Rightarrow f(\sigma) \geq r - 1$$

## Proposition:

- For any marking algorithm, the number of misses it incurs on any sequence  $\sigma$  is at most  $k \cdot f(\sigma) + k$

## Randomized Caching

---

### Algorithm 2: A Randomized Marking Algorithm

---

Each memory item can be either **marked** or **unmarked**

At the beginning of the phase, all items are unmarked

On a request to item  $s$ :

Mark  $s$

**if**  $s$  *is in the cache* **then**

    | Evict nothing

**else**

**if** *All items currently in the cache are marked* **then**

        | Declare the phase over

        | Processing of  $s$  is deferred to start of next phase

**else**

        | **Evict** an **unmarked item chosen uniformly at random** from  
        | the cache

**end**

**end**

---

## Chernoff Bounds

## Reminder

- We defined earlier the expectation of a random variable and worked with this definition
- Intuitively, we have a sense that the value of a random variable ought to be “near” its expectation with reasonably high probability, but we have not yet explored the extent to which this is true

## Definition:

- We say that two random variables  $X$  and  $Y$  are **independent** if, for any values  $i$  and  $j$ , the events  $\Pr[X = i]$  and  $\Pr[Y = j]$  are independent
- This definition extends naturally to larger sets of random variables



## Chernoff Bounds

- Consider a random variable  $X$  that is a sum of several independent 0-1-valued random variables:  $X = X_1 + X_2 + \dots + X_n$ , where  $X_i$  takes value 1 with probability  $p_i$  and the value 0 otherwise
- By the linearity of the expectation, we have:

$$E[X] = \sum_{i=1}^n p_i$$

- Intuitively, the independence of the random variables  $X_1, \dots, X_n$  suggests that their fluctuations are likely to “cancel out”, and so their sum  $X$  will have a value close to its expectation with high probability

# Chernoff Bounds

## First bound:

- We bound the probability that  $X$  deviates above  $E[X]$

## Second bound:

- We bound the probability that  $X$  deviates below  $E[X]$

These two results are called the Chernoff bounds

## Chernoff Bounds

### Theorem:

Let  $X, X_1, X_2, \dots, X_n$  be defined as above, and assume that  $\mu \geq E[X]$ . Then for any  $\delta > 0$ , we have:

$$\Pr[X > (1 + \delta)\mu] < \left[ \frac{e^\delta}{(1 + \delta)^{(1+\delta)}} \right]^\mu$$

This means that sum of independent 0-1 random variables is tightly centered on the mean

### Proof.

We apply a number of simple transformations

- For any  $t > 0$ , we have:

$$\Pr[X > (1 + \delta)\mu] = \Pr[e^{tX} > e^{t(1+\delta)\mu}] \leq e^{-t(1+\delta)\mu} E[e^{tX}]$$

- The first equality derives from the fact that  $e^{tX}$  is monotone in  $x$
- The second inequality derives from the Markov's inequality  $\Pr[X > a] \leq E[X]/a$

□

# Chernoff Bounds

## Proof.

- Now,  $E[e^{tX}] = E[e^{t \sum_i X_i}] = \prod_i E[e^{tX_i}]$
- First equality is the definition of  $X$
- Second equality is due to independence



# Chernoff Bounds

## Proof.

- Let  $p_i = \Pr[X_i = 1]$ . Then,

$$E[e^{tX_i}] = p_i e^t + (1 - p_i)e^0 = 1 + p_i(e^t - 1) \leq e^{p_i(e^t - 1)}$$

- Indeed, for any  $\alpha \geq 0$ ,  $1 + \alpha \leq e^\alpha$
- Combining everything we have:

$$\Pr[X > (1 + \delta)\mu] \leq e^{-t(1+\delta)\mu} \prod_i E[e^{tX_i}] \leq e^{-t(1+\delta)\mu} \prod_i e^{p_i(e^t - 1)} \leq e^{-t(1+\delta)\mu} e^{\mu(e^t - 1)}$$

- And finally, we must choose  $t = \ln 1 + \delta$



# Chernoff Bounds

## Theorem:

Let  $X, X_1, X_2, \dots, X_n$  be defined as above, and assume that  $\mu \geq E[X]$ . Then for any  $\mu \leq E[X]$  and for any  $0 \leq \delta \leq 1$ , we have:

$$\Pr[X < (1 - \delta)\mu] < e^{-\delta^2 \mu / 2}$$

- The proof is similar to that we've seen before

## Randomized On-line Load Balancing

## The problem

- Suppose we have a system in which  $m$  jobs arrive in a stream and need to be processed immediately
- We have a collection of  $n$  identical processors that are capable of performing the jobs
- The goal is to assign each job to a processor in a way that balances the workload evenly across the processors



# Load Balancing

## The challenge

- Assume that the systems lacks the coordination or centralization to implement what described before
- A lightweight approach would be to simply assign each job to one of the processors **uniformly at random**
- Intuitively, this should also balance the jobs evenly, since each processor is equally likely to get each job
- At the same time, since the assignment is completely random, one doesn't expect everything to end up perfectly balanced

How well does the randomized algorithm behave?

## Further notes

- This is similar to what has been discussed for hash functions
  - There, instead of assigning jobs to processors, we were assigning elements to entries in a hash table
- ⇒ The analysis we do in this part, is also relevant to the study of hashing schemes

## Analyzing a Random Allocation

- We will see that the analysis of our random load balancing process depends on the relative size of  $m$ , the number of jobs, and  $n$ , the number of processors
  - We start with a particular case:  $m = n$
- ⇒ In this case it is possible for each processor to end up with exactly 1 job assigned, although this is not very likely
- Instead, we expect that some processors will receive no jobs, and others will receive more than one job
- ⇒ We study how heavily loaded with jobs a processor can become

## Load Balancing

### The case $m = n$

- Let  $X_i$  be the random variable equal to the number of jobs assigned to processor  $i$ , for  $i = 1, 2, \dots, n$ .
- It is easy to determine the expected value of  $X_i$
- Let  $Y_{ij}$  be the random variable equal to 1 if job  $j$  is assigned to processor  $i$  and 0 otherwise
- Then we have:

$$X_i = \sum_{j=1}^n Y_{ij}$$

$$E[Y_{ij}] = \frac{1}{n}$$

$$\Rightarrow E[X_i] = \sum_{j=1}^n E[Y_{ij}] = 1$$

## Load Balancing

- Our concern is with how far  $X_i$  can deviate above its expectation

What is the probability that  $X_i > c$ ?

- To give an upper bound on this, we can directly apply the Chernoff bound: indeed  $X_i$  is the sum of independent 0-1-valued random variables  $Y_{ij}$ , where  $\mu = 1$  and  $1 + \delta = c$

What is the probability that  $X_i > c$ ?

### Proposition

- When  $m = n$  we have:

$$\Pr[X_i > c] < \left( \frac{e^c - 1}{c^c} \right)$$

- In order for there to be a small probability for **any**  $X_i$  exceeding  $c$ , we will take the Union Bound over  $i = 1, \dots, n$
- So we need to choose  $c$  large enough to drive  $\Pr[X_i > c]$  well below  $1/n$  for each  $i$

## Load Balancing

- This requires looking at the denominator of the Chernoff Bound
  - We need to understand how  $c^c$  grows with  $c$ , and make it large enough
- ⇒ We need to study what is the  $x$  such that  $x^x = n$
- Suppose we let  $\gamma(n)$  be this number  $x$
  - There is no closed-form expression for  $\gamma(n)$ , but we can determine its asymptotic value as follows
  - If  $x^x = n$ , then taking logarithms we have  $x \log(x) = \log(n)$
  - And taking logarithms again we have  $\log(x) + \log \log(x) = \log \log(n)$

$\Rightarrow 2 \log x > \log x + \log \log x = \log \log n > \log x$

- We use this to divide through the equation  $x \log x = \log n$ :

$$\frac{1}{2}x \leq \frac{\log n}{\log \log n} \leq x = \gamma(n)$$

- Thus:

$$\gamma(n) = \Theta\left(\frac{\log n}{\log \log n}\right)$$



## Load Balancing

- Now, if we set  $c = e^{\gamma(n)}$ , then we have:

$$\begin{aligned}\Pr[X_i > c] &< \left(\frac{e^c - 1}{c^c}\right) < \left(\frac{e}{c}\right)^c = \\ &= \left(\frac{1}{\gamma(n)}\right)^{e^{\gamma(n)}} < \left(\frac{1}{\gamma(n)}\right)^{2\gamma(n)} = \frac{1}{n^2}\end{aligned}$$

- Thus, applying the Union Bound for  $X_1, \dots, X_n$ , we get the following

### Theorem:

With probability at least  $\left(1 - \frac{1}{n}\right)$ , no processor receives more jobs than:

$$e^{\gamma(n)} = \Theta\left(\frac{\log n}{\log \log n}\right)$$

- With a more involved analysis, one can also show that this bound is **asymptotically tight**: with high probability, some processor actually receives a number of jobs bounded by:

$$\Omega\left(\frac{\log n}{\log \log n}\right)$$

Increasing the number of jobs (**i.e.**, going beyond  $m = n$ )

- We now use Chernoff bounds to argue that, as more jobs are introduced into the system, the loads “smooth out” rapidly
- ⇒ The number of jobs on each processor quickly becomes equalized within some constant factors

## Load Balancing

Assume  $m = 16n \ln n$  jobs

- The expected load per processor is  $\mu = 16 \ln n$
- Using the first Chernoff Bound, we see that the probability of any processor's load exceeding  $32 \ln n$  is at most:

$$\Pr[X_i > 2\mu] < \left(\frac{e}{4}\right)^{16 \ln n} < \left(\frac{1}{e^2}\right)^{\ln n} = \frac{1}{n^2}$$

- Also, the probability that any processor's load is below  $8 \ln n$  is:

$$\Pr[X_i < \frac{1}{2}\mu] < e^{-\frac{1}{2}\left(\frac{1}{2}\right)^2 (16 \ln n)} = e^{-2 \ln n} = \frac{1}{n^2}$$

# Load Balancing

## Theorem:

When there are  $n$  processors and  $\Omega(n \log n)$  jobs, then with high probability, every processor will have a load between **half** and **twice** the average